Franklin W. Olin College of Engineering

MTH3199 Applied Math for Engineers – Fall 2025
Assignment 3: Introduction to Numerical Integration

Assigned: Friday, October 3, 2025.
Lab Report Due: Thursday, October 16, 2025 (11:59 PM EST).

Online Resources

- Forward Euler: Wikipedia

- Backward Euler: Wikipedia

- Explicit & Implicit midpoint method: Wikipedia

- Cool youtube video comparing various methods

Overview

In today's activity, we will begin to look at a few different numerical integration techniques. Our goal is to find an approximate solution to the following initial value problem (IVP):

$$\dot{X} = f(t, X), \qquad X(t_{start}) = X_0 \tag{1}$$

In this differential equation, the independent variable, t, is a scalar, and the independent variable, X, can either be a scalar or a vector. The function $f(t, X)$ describes the rate at which X changes over time. The solution to the IVP is a function of time, $X(t)$, that satisfies the IC's, $X(t_{start}) = X_0$, and the differential equation, $\dot{X} = f(t, X)$.

Broadly speaking, there are two categories of numerical methods for solving ODE's:

- **Explicit methods** directly compute the system state at the next time step, X_{n+1}, as a function of the the current state, X_n.

- **Implicit methods** set up a (possibly nonlinear) set of equations that constrain the state at the next time step, X_{n+1}. These equations must then be solved (often with Newton's method) to find X_{n+1}.

To get a feel for how numerical integration works, we will start by implementing and testing two basic explicit methods: forward Euler and explicit midpoint:

- Forward Euler approximates the value of X_{n+1} at the next time step by using the tangent line approximation.

$$X(t + h) \approx X(t) + h\dot{X} \quad \rightarrow \quad X_{n+1} = X_n + hf(t_n, X_n) \tag{2}$$

  where h is the length of the time step.

- The explicit midpoint method uses a forward Euler step to approximate the value of $X_{n+.5}$ a half-time step away (i.e. the **midpoint** between X_n and X_{n+1}). It then evaluates the slope at this midpoint to approximate X_{n+1} at the next time step.

$$X_{n+.5} = X_n + \frac{h}{2}f(t_n, X_n) \tag{3}$$

$$X_{n+1} = X_n + hf(t_n + \frac{h}{2}, X_{n+.5}) \tag{4}$$

  where h is the length of the time step. It should be noted that this method is different than just computing two forward Euler steps using a time step of $\frac{h}{2}$.

One last thing before we get into the implementation section: using computers to solve differential equations is an entire field of applied mathematics. There's a whole zoo of integration algorithms out there, and the content presented in this module is just the tip of the iceberg.

Instructions

Before you begin, you are encouraged to form groups of two or three. Groups will jointly submit a single set of deliverables on Canvas (and will be graded as a single unit). Once you have found teammates (or if you are choosing to work alone), add yourself (and your teammates if you have them) to the same **assignment 03** group on Canvas.

Test Functions

To test the performance of our numerical integration scheme, we'll consider two initial value problems (IVP's) that have known analytical/closed-form solutions. The first IVP describes a first order system driven by a sinusoid:

$$\dot{x} = -5x + 5\cos(t) - \sin(t), \quad x(0) = 1 \tag{5}$$

The solution to this IVP is $x(t) = \cos(t)$. I have provided an implementation for both the rate function, $f(t, x)$, and the corresponding solution, $x(t)$, below:

```
function dXdt = rate_func01(t,X)
    dXdt = -5*X + 5*cos(t) - sin(t);
end

function X = solution01(t)
    X = cos(t);
end
```

The second IVP describes an undamped second order system:

$$\dot{X} = \begin{bmatrix} 0, & -1 \\ 1, & 0 \end{bmatrix} X, \quad X(0) = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \tag{6}$$

and has the solution:

$$X(t) = \begin{bmatrix} \cos(t) \\ \sin(t) \end{bmatrix} \tag{7}$$

I have provided an implementation for both the rate function, $f(t, X)$, and solution, $X(t)$, below:

```
function dXdt = rate_func02(t,X)
    dXdt = [0,-1;1,0]*X;
end

function X = solution02(t)
    X = [cos(t);sin(t)];
end
```

Implementing Forward Euler

Your first task is to implement forward Euler:

$$X_{n+1} = X_n + hf(t_n, X_n) \tag{8}$$

We'll do this in two steps. First, write a MATLAB function that computes X_{n+1} for the next time step, given the rate function $f(t, X)$ the current values of t_n and X_n, and the size of the time step, h. For the purpose of performance experiments later in this activity, this function should also return the number of times it called the rate function, $f(t, X)$. I have provided a template below:

```
%This function computes the value of X at the next time step
%using the Forward Euler approximation
```

```
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)   (t is before X)
%t: the value of time at the current step
%XA: the value of X(t)
%h: the time increment for a single step i.e. delta_t = t_{n+1} - t_{n}
%OUTPUTS:
%XB: the approximate value for X(t+h) (the next step)
%    formula depends on the integration method used
%num_evals: A count of the number of times that you called
%           rate_func_in when computing the next step
function [XB,num_evals] = forward_euler_step(rate_func_in,t,XA,h)
    %your code here
end
```

Next, use your single step function to write a MATLAB function that runs the numerical integration over some time interval, $t \in [t_0, t_f]$, given the rate function $f(t, X)$, the initial condition X_0, and some target step size, h_{ref}. Note that, since we would like the step sizes to be equal, it won't usually be possible to set the step size to be exactly h_{ref}. Instead, just find the smallest **number of steps**, N, such that the step size $h = \frac{t_f - t_0}{N}$ is still smaller than h_{ref}. This function should keep track of the number of times the rate function was called.

```
%Runs numerical integration using forward Euler approximation
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)   (t is before X)
%tspan: a two element vector [t_start,t_end] that denotes the integration endpoints
%X0: the vector describing the initial conditions, X(t_start)
%h_ref: the desired value of the average step size (not the actual value)
%OUTPUTS:
%t_list: the vector of times, [t_start;t_1;t_2;...;.t_end] that X is approximated at
%X_list: the vector of X, [X0';X1';X2';...;(X_end)'] at each time step
%h_avg: the average step size
%num_evals: total number of calls made to rate_func_in during the integration
function [t_list,X_list,h_avg, num_evals] = ...
            forward_euler_fixed_step_integration(rate_func_in,tspan,X0,h_ref)
    %your code here
end
```

To verify that your implementation is working, use it to solve the IVP described by equation 5. Plot the closed-form solution for $X(t)$ and a few numerical approximations (with different time steps) on the same axes.

Explicit Midpoint Method

Your next task is to implement the explicit midpoint method:

$$X_{n+.5} = X_n + \frac{h}{2}f(t_n, X_n) \tag{9}$$

$$X_{n+1} = X_n + hf(t_n + \frac{h}{2}, X_{n+.5}) \tag{10}$$

As with forward Euler, we'll do this in two steps. First, write a MATLAB function that computes X_{n+1} for the next time step, given the rate function $f(t, X)$ the current values of t_n and X_n, and the size of the time step, h. For the purpose of performance experiments later in this activity, this function should also return the number of times it called the rate function, $f(t, X)$. I have provided a template below:

```
%This function computes the value of X at the next time step
%using the explicit midpoint approximation
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)   (t is before X)
%t: the value of time at the current step
```

```
%XA: the value of X(t)
%h: the time increment for a single step i.e. delta_t = t_{n+1} - t_{n}
%OUTPUTS:
%XB: the approximate value for X(t+h) (the next step)
%    formula depends on the integration method used
%num_evals: A count of the number of times that you called
%           rate_func_in when computing the next step
function [XB,num_evals] = explicit_midpoint_step(rate_func_in,t,XA,h)
    %your code here
end
```

Next, use your single step function to write a MATLAB function that runs the numerical integration over some time interval, $t \in [t_0, t_f]$, given the rate function $f(t, X)$, the initial condition X_0, and some target step size, h_{ref}. Note that, since we would like the step sizes to be equal, it won't usually be possible to set the step size to be exactly h_{ref}. Instead, just find the smallest **number of steps**, N, such that the step size $h = \frac{t_f - t_0}{N}$ is still smaller than h_{ref}. This function should keep track of the number of times the rate function was called.

```
%Runs numerical integration using explicit midpoint approximation
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)  (t is before X)
%tspan: a two element vector [t_start,t_end] that denotes the integration endpoints
%X0: the vector describing the initial conditions, X(t_start)
%h_ref: the desired value of the average step size (not the actual value)
%OUTPUTS:
%t_list: the vector of times, [t_start;t_1;t_2;...;.t_end] that X is approximated at
%X_list: the vector of X, [X0';X1';X2';...;(X_end)'] at each time step
%h_avg: the average step size
%num_evals: total number of calls made to rate_func_in during the integration
function [t_list,X_list,h_avg, num_evals] = ...
            explicit_midpoint_fixed_step_integration(rate_func_in,tspan,X0,h_ref)
    %your code here
end
```

To verify that your implementation is working, use it to solve the IVP described by equation 5. Plot the closed-form solution for $X(t)$ and a few numerical approximations (with different time steps) on the same axes.

Experiment: Local Truncation Error

Truncation error is the error that originates from the discrete nature of the numerical integration method (the number of time steps are finite). This is different from round-off error which is a consequence of computers not being able to represent real numbers with infinite precision. There are two types of truncation error:

- **Local:** the error for a single time step.

- **Global:** the error across the entire integration interval.

Decreasing the size of the time step will cause both of these truncation errors to shrink. Unfortunately, smaller time steps means more steps, which means more computation time. The rate at which the truncation error, ϵ scales with the time step, h varies from method to method. We can describe this relationship using big O notation:

$$\epsilon = O(h^p) \tag{11}$$

Big O notation is used to describe the asymptotic behavior of a function. In this case, as $h \to 0$, the error, ϵ is become proportional to h^p:

$$\epsilon \approx kh^p \tag{12}$$

The proportionality constant, k, is not relevant to big O notation, all that matters is that ϵ and h^p are proportional to one another when h is sufficiently small.

Your task is to run a set of experiments to measure how the local truncation error scales with the size of the time step (in other words, find p). It should be noted that the theoretical value of p is a whole number. Let $X_{n+1} =$

$G(t, X_n, h)$ be the function that computes next value of X_{n+1}. If $X(t)$ is the analytical solution to the IVP, then the local truncation error is given by:

$$\epsilon = |G(t, X(t), h) - X(t+h)| \tag{13}$$

Note that $|X|$ is the norm of X when X is a vector. Choose one of the test functions and some reference time t_{ref} (**note**: do not choose $t_{ref} = 0$ or $t_{ref} = \pi n$! just choose an arbitrary value, like $t_{ref} = .492$). For each algorithm, compute the local truncation error at $t = t_{ref}$ for different values of h (I'd like to see around 100 data points ranging from 10^{-5} to 10^1 that are distributed evenly in log-space). Fit a line to this data in log space to estimate the value of p. I have uploaded a MATLAB function to Canvas (called log-log regression) to compute the fit for you. On the same axes, plot the data you gathered for both algorithms, as well as the corresponding fit lines (this should be a log-log plot). What were the estimated values of p that you computed for each algorithm?

Part of the reason why the local truncation error is shrinking is that both $X(t+h)$ and $G(t, X(t), h)$ are approaching $X(t)$ as the step size, h, gets smaller and smaller. To verify that the local truncation error is shrinking in a meaningful way (and not just for the aforementioned reason), plot the difference $|X(t+h) - X(t)|$ as a function of h (for the same values of h that you already used) on the same axes, as well as a fit line.

Run the same experiment using the other test function. You do not need to generate additional plots. However, please include the new estimates of p in your lab report. Based on your results, does the local truncation error shrink in a meaningful way (does it shrink faster than $|X(t+h) - X(t)|$)? Does it shrink faster for the forward Euler step or the explicit midpoint step?

Experiment: Global Truncation Error

Now, let's see how the global truncation error scales with the size of the time step, h. If we use our numerical integration scheme to approximate X over the interval $t \in [t_0, t_f]$ with initial condition $X(t_0) = X_0$, and $X(t)$ is the closed-form solution, then the global truncation error is given by:

$$\epsilon = |X_f - X(t_f)| \tag{14}$$

where X_f is the numerical approximation of $X(t)$ after the last time step. Choose one of the test functions and some time interval $t \in [t_0, t_f]$. For each algorithm, compute the global truncation error for different values of h (I would like to see values that are evenly spaced across several orders of magnitude). Fit a line to this data in log space to estimate the value of p. On the same axes, plot the data you gathered for both algorithms, as well as the corresponding fit lines (this should be a log-log plot). What were the estimated values of p that you computed for each algorithm?

Your implementations should keep track of how many times the rate function $f(t, X)$ was called. How does the global truncation error scale with the number of rate function calls (i.e. what are the corresponding values of p)? On a separate figure, generate a plot comparing this scaling for the two algorithms (this plot should include both the experimental data and their fit lines).

Run the same set of experiments using the other test function. You do not need to generate additional plots. However, in your lab report, please include the new estimates of p indicating how the global truncation error scales with both the step size, and the number of rate function calls.

Implicit Methods

Day 11 (Tuesday, October 7th) activity. Please complete before class on Friday, October 10th.

Now that we've implemented two explicit methods, let's try our hand at generating the equivalent implicit methods. As previously mentioned, implicit methods set up a system of equations that constrain the state at the next time step, X_{n+1}, which must then be solved. Implicit methods can often be more computationally expensive (since solving a system of nonlinear equations can take some effort). However, implicit methods can guarantee that the integration will be stable (more on this later). For the moment, let's look at two examples, backward Euler and implicit midpoint:

- Like forward Euler, backward Euler uses a tangent line approximation to relate X_n and X_{n+1}. The key difference is that, for backward Euler, the tangent line starts at (t_{n+1}, X_{n+1}):

$$X(t) \approx X(t+h) - h\dot{X}(t+h) \quad \rightarrow \quad X_n = X_{n+1} - hf(t_n + h, X_{n+1}) \tag{15}$$

Moving everything to one side, we get:

$$G(X_{n+1}) = X_n + hf(t_n + h, X_{n+1}) - X_{n+1} = 0 \qquad (16)$$

In other words, X_{n+1} is defined implicitly (it's the root of function G). Depending on the rate function, $f(t, X)$, it may be possible to write a formula for X_{n+1}. However, most of the time we will need to use an algorithm like Newton's method (wink wink) to solve for X_{n+1}.

- Implicit midpoint is very similar to explicit midpoint, except that the midpoint, $X_{n+.5}$, is the actual average of X_n and X_{n+1}, instead of being approximated through a forward Euler step:

$$X_{n+.5} = \frac{1}{2}X_n + \frac{1}{2}X_{n+1} \qquad (17)$$

$$X_{n+1} = X_n + hf(t_n + \frac{h}{2}, X_{n+.5}) \qquad (18)$$

This can be rewritten as the following implicit equation:

$$G(X_{n+1}) = X_n + hf\left(t_n + \frac{h}{2}, \frac{1}{2}(X_n + X_{n+1})\right) - X_{n+1} \qquad (19)$$

which you can solve using Newton's method.

Implementing Backward Euler

Your task is to implement backward Euler:

$$G(X_{n+1}) = X_n + hf(t_n + h, X_{n+1}) - X_{n+1} = 0 \qquad (20)$$

We'll do this in two steps. First, write a MATLAB function that computes X_{n+1} for the next time step, given the rate function $f(t, X)$ the current values of t_n and X_n, and the size of the time step, h. For the purpose of performance experiments, this function should also return the number of times it called the rate function, $f(t, X)$. To do this, you will need to first construct the function $G(X_{n+1})$, and then use your multidimensional Newton solver to solve for X_{n+1}. I'd recommend that you use X_n as your initial guess for the solver. Since we want to count the number of calls made to the rate function, you will need to edit your Newton solver and your Jacobian approximation function so that they keep track of the total number of function calls. I have provided a template below for you to fill out:

```
%This function computes the value of X at the next time step
%using the Backward Euler approximation
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)   (t is before X)
%t: the value of time at the current step
%XA: the value of X(t)
%h: the time increment for a single step i.e. delta_t = t_{n+1} - t_{n}
%OUTPUTS:
%XB: the approximate value for X(t+h) (the next step)
%    formula depends on the integration method used
%num_evals: A count of the number of times that you called
%           rate_func_in when computing the next step
function [XB,num_evals] = backward_euler_step(rate_func_in,t,XA,h)
    %your code here
end
```

Next, use your single step function to write a MATLAB function that runs the numerical integration using equal time steps over some time interval, $t \in [t_0, t_f]$, given the rate function $f(t, X)$, the initial condition X_0, and some target step size, h_{ref}. You may have noticed that the code that runs the fixed-step integration for the different methods is basically identical, with the only difference being which approximation method is called at each step. Thus, let's just write a more general version that takes the approximation function as a input, and eliminate the redundant code:

```
%Runs fixed step numerical integration
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)   (t is before X)
%step_func: the function used to approximate X(t) at the next step
%              [XB,num_evals] = step_func(rate_func_in,t,XA,h)
%tspan: a two element vector [t_start,t_end] that denotes the integration endpoints
%X0: the vector describing the initial conditions, X(t_start)
%h_ref: the desired value of the average step size (not the actual value)
%OUTPUTS:
%t_list: the vector of times, [t_start;t_1;t_2;...;.t_end] that X is approximated at
%X_list: the vector of X, [X0';X1';X2';...;(X_end)'] at each time step
%h_avg: the average step size
%num_evals: total number of calls made to rate_func_in during the integration
function [t_list,X_list,h_avg, num_evals] = ...
            fixed_step_integration(rate_func_in,step_func,tspan,X0,h_ref)
    %your code here
end
```

To verify that your implementation is working, use it to solve the IVP described by equation 5. Plot the closed-form solution for $X(t)$ and a few approximations (with different time steps) on the same axes.

Implicit Midpoint Method

Your next task is to implement the implicit midpoint method:

$$G(X_{n+1}) = X_n + hf\left(t_n + \frac{h}{2}, \frac{1}{2}(X_n + X_{n+1})\right) - X_{n+1} \tag{21}$$

Once again, we'll do this in two steps. First, write a MATLAB function that computes X_{n+1} for the next time step, given the rate function $f(t, X)$ the current values of t_n and X_n, and the size of the time step, h. For the purpose of performance experiments later in this activity, this function should also return the number of times it called the rate function, $f(t, X)$. Like backward Euler, you will need to construct $G(X_{n+1})$, and then use your multidimensional Newton solver to compute X_{n+1}:

```
%This function computes the value of X at the next time step
%using the implicit midpoint approximation
%INPUTS:
%rate_func_in: the function used to compute dXdt. rate_func_in will
%              have the form: dXdt = rate_func_in(t,X)   (t is before X)
%t: the value of time at the current step
%XA: the value of X(t)
%h: the time increment for a single step i.e. delta_t = t_{n+1} - t_{n}
%OUTPUTS:
%XB: the approximate value for X(t+h) (the next step)
%    formula depends on the integration method used
%num_evals: A count of the number of times that you called
%           rate_func_in when computing the next step
function [XB,num_evals] = implicit_midpoint_step(rate_func_in,t,XA,h)
    %your code here
end
```

Use you generalized fixed-step integration function to solve the IVP described by equation 5 using the implicit midpoint method. Plot the closed-form solution for $X(t)$ and a few numerical approximations (with different time steps) on the same axes.

Explicit vs. Implicit: Truncation Error

Generate the following plots and add the following entries to your tables:

- Create a single log-log plot that compares the **local** truncation error (as a function of step size) for forward/backward Euler and explicit/implicit midpoint (for the IVP described by eqn. 5). **Do not include fit**

lines! (too much clutter). **Do not include the analytical difference** $|X(t + h) - X(t)|$ (too much clutter). Remember to use a nonzero value for t_{ref}.

- Create a single log-log plot that compares the **global** truncation error (as a function of step size) for forward/backward Euler and explicit/implicit midpoint (for the IVP described by eqn. 5). **Do not include fit lines!** (too much clutter).

- Create a single log-log plot that compares the **global** truncation error scaling (with the **number of rate function calls**) (for the IVP described by eqn. 5). **Do not include fit lines!** (too much clutter).

- Add entries to all of your tables for backward Euler and implicit midpoint.

Based on the results of these experiments, are implicit methods an improvement over their analogous explicit methods? Why or why not?

Explicit vs. Implicit: Stability

Let's take a closer look at how our numerical solutions grow/shrink over time:

- Create a figure that compares the numerical solutions to eqn. 5 generated using forward/backward Euler and explicit/implicit midpoint, with a target time step of $h_{ref} = .38$.

    - This should be a plot of the solution $x(t)$ as a function of time.
    - Try to choose a time interval around $[t_0, t_f] = [0, 20]$.
    - I'd recommend that you format this figure as four separate subplots (on top of one another), each showing one of the numerical solutions juxtaposed with the analytical solution.

- Now, perform the same exact experiment, with a target time step of $h_{ref} = .45$.

    - Generate a separate figure for this value of h_{ref}.

- Based on the experimental results, are implicit methods an improvement over explicit methods?

    - Why or why not?

Deliverables and Submission Guidelines

This exercise is the first part of the larger numerical integration assignment. The lab report for the numerical integration assignment will be due on Thursday, October 16th at 11:59 PM EST. In this lab report, make sure to document how you went about implementing the forward Euler and the explicit midpoint method. As a reminder, I want to see the following plots and tables:

- Plots comparing the closed-form solution to equation 5 with the numerical approximation (for a few different time step sizes). The idea is that each plot should depict that numerical approximation of $x(t)$ getting more and more accurate as the step size gets smaller and smaller.

  - Each algorithm should get its own plot.

- Plots comparing the numerical solutions to eqn. 5 for a time step of $h_{ref} = .38$ and a separate plot for $h_{ref} = .45$ (as detailed in the explicit vs. implicit stability section).

  - Based on the experimental results, are implicit methods an improvement over explicit methods?
  - Why or why not?

- A table that shows how the **local** truncation error scales with step size (the values of p). This table should have entries for each algorithm/test function combination.

Method	Estimate for p (test 1)	Estimate for p (test 2)
$\|X(t+h) - X(t)\|$		
Forward Euler		
Backward Euler		
Explicit Midpoint		
Implicit Midpoint		

- A log-log plot that compares the **local** truncation error scaling for the two **explicit** methods.

  - This plot should include the difference $|X(t+h) - X(t)|$ for comparison.
  - Step sizes should be across a range of magnitudes.
  - The plot should include both the data points and fit lines.
  - You only need to generate a plot for a single test function.

- A log-log plot that compares the **local** truncation error scaling (as a function of step size) for all four methods (forward/backward Euler and explicit/implicit midpoint).

  - **Do not include fit lines or the analytical difference** $|X(t+h) - X(t)|$.
  - Remember to use a nonzero value for t_{ref}.

- A table that shows how the **global** truncation error scales with step size (the values of p). This table should have entries for each algorithm/test function combination.

Method	Estimate for p (test 1)	Estimate for p (test 2)
Forward Euler		
Backward Euler		
Explicit Midpoint		
Implicit Midpoint		

- A short discussion on which local truncation error scaling is preferable (forward Euler or explicit midpoint).

- A log-log plot that compares the **global** truncation error scaling for the two **explicit** methods. Step sizes should be across a range of magnitudes. Please include both the data points and fit lines. You only need to generate a plot for a single test function.

- A log-log plot that compares the **global** truncation error scaling (as a function of step size) for all four methods (forward/backward Euler and explicit/implicit midpoint). **Do not include fit lines**.

- A table that shows how the **global** truncation error scales with the **number of rate function calls** (the values of p). This table should have entries for each algorithm/test function combination.

Method	Estimate for p (test 1)	Estimate for p (test 2)
Forward Euler		
Backward Euler		
Explicit Midpoint		
Implicit Midpoint		

- A log-log plot that compares the **global** truncation error scaling (with the **number of rate function calls**) for the two **explicit** methods (Forward Euler and Explicit Midpoint).

  - Please include both the data points and fit lines.
  - You only need to generate a plot for a single test function.

- A log-log plot that compares the **global** truncation error scaling (with the **number of rate function calls**) for all four methods (forward/backward Euler and explicit/implicit midpoint). **Do not include fit lines**.

- The explicit vs. implicit stability visualizations.

- A short discussion on which algorithm you would prefer to use for integrating these test functions and why.

- Please make sure to include some interpretation for what each of the plots tell you. You don't need a blurb for every single plot, but there should be at least some discussion about what you think these results mean.